

$$\text{ICL} = \text{GD}$$

Mathematics of Generative Models, Spring 2024

# Proposal

In-context learning forms the basis of modern models' capabilities.

As a behaviour it is poorly understood.

Various formulations have been proposed. This work uses gradient descent.

We extend this to a testbed of formal languages.

Extend to a meaningful toy setting – formal languages.

Well-studied and easily controllable.

Better than the random task used in the paper – more interpretable insights.

Each input sample consists of a set samples  $(x_i, y_i)$ .

This is followed by a test sample  $x_{\text{test}}$ .

For nonlinear tasks, the embeddings are passed through an MLP first.

An LSA layer is trained on these representations to solve the task.

## Extension 1: Domain

Use a string membership task in a fixed formal language.

Formulate as regression by fitting to  $\pm 1$ .

We choose regular languages.

## Extension 2: Setup

The authors train an end-to-end model for nonlinear regressions. This uses an MLP and an LSA layer.

One issue with this is that the separation between linearization and GD is presupposed (mid-eval).

We fix this by using an independently trained autoencoder.

# Experiments

We use the following regular languages:

$$a * bc*; ((a|b)c)*; (ab[c])*$$

First, we train an autoencoder to encode strings from these languages.

We then train an LSA layer to predict in-context string membership from this representation.

We also construct an LSA layer that is expected to perform GD on the inputs.

Throughout training, we compare the predictions of the training model with the predictions of the construction (MSE).

We also compare heatmaps of the learned LSA layer with the construction.



# Results: Autoencoder Validation Loss




Figure 1: Autoencoder Validation Loss

# Results: Classifier Validation Accuracy and Loss



Figure 2: Classifier Validation Accuracy



Figure 3: Classifier Validation Loss

# Results: MSE of Classifier Predictions




Figure 4: MSE between Predictions

# Results: $QK$ Heatmaps




Figure 5: Constructed MHA Heatmap




Figure 6: Learned MHA Heatmap

# Results: *PV* Heatmaps




Figure 7: Constructed MHA Heatmap




Figure 8: Learned MHA Heatmap

The model learns to solve the task with an accuracy of about 75%.

The results do not match the authors' predictions.

There are multiple possibilities for this:

- The authors' setup relies on the MLP-LSA end-to-end training (extension 2), which we do away with.
- The task is “too” non-linear.